



CHINA LINUX KERNEL  
中国Linux内核开发者大会

# PMR：并行内存回收的设计与实践

李文通  
OPPO高级底层软件工程师  
2025年11月1日



# 目 录

CONTENTS

01

## 背景

Background

02

## 动机

Motivation

03

## 设计

PMR's Design

04

## 实验与总结

Evaluation and Conclusion




# Today's Mobile Devices are facing Memory Pressure

Memory capacity is becoming a scarce resource on mobile devices

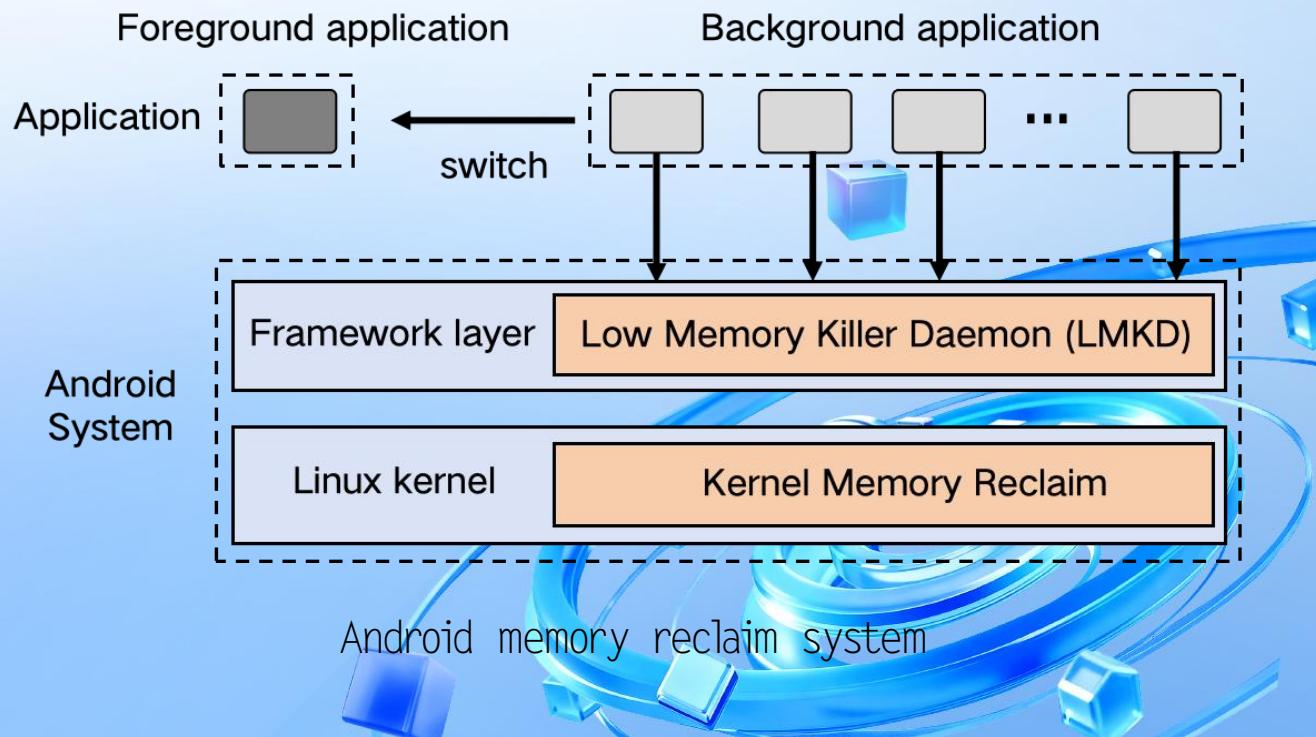
The application memory footprint have been growing.<sup>[1]</sup>

The physical memory capacity of mobile devices has seen only modest increase.

Mobile systems need efficient memory reclaim to relieve memory pressure





Application memory footprint grows rapidly



# Memory Reclaim on Mobile device

Kernel memory reclamation is more cost-efficient than LMKD


It includes memory swapping and direct reclaim, and both of them follow the same reclaim path: page shrinking and page writeback




# Preliminary Study of System Memory Reclaim

## Frequent, Expensive Process Killing

LMKD is frequently triggered, causing the application to cold launch, accompanied by a longer response time.



Many processes are killed among different mobile devices





# Preliminary Study of System Memory Reclaim

## Sluggish Kernel Memory Reclaim

With the upgrade of mobile devices, the memory reclaim throughput has grown slowly.

The maximum throughput is only over 100 MB, which is far lower than the memory requirements of the popular applications.





While app memory footprint is increasingly high, the hardware advance does not benefit the efficiency of kernel memory reclaim.

# Performance Bottleneck Analysis

## I/O Utilization

Low Impact of I/O Conflicts

Excessive Storage Bandwidth




I/O Performance is not the bottleneck of kernel memory reclaim.

# Performance Bottleneck Analysis

## Sub-optimal Memory Reclaim Path

**Sequential Execution of Key Steps:** page writeback waits on the results of page shrinking, introducing unnecessary delays.




# Performance Bottleneck Analysis


## Sub-optimal Memory Reclaim Path

**Inefficient Page Shrinking:** page shrinking suffers from inefficient reclaim and repeated invocation.

**Internally Blocked Page Writeback:** page unmap latency is highly unstable and produces small writes that cannot unleash the potential of flash storage.



nr\_to\_reclaim vs nr\_reclaimed when system performs memory reclaim




# PMR's Design

PMR parallelizes key steps of memory reclaim to fulfill application memory demand

Proactive Page Shrinking(PPS)

Storage-friendly Page Writeback(SPW)



# PMR's Design

## Proactive Page Shrinking(PPS)

### Decoupling Page Shrinking and Page Writeback:

- (1) Enable an independent *kshrinkd* Thread.
- (2) Maintain a victim page list.

### Replenishing Victim Pages:

- (1) Design page shrinking watermark: The parameter *nr\_ideal\_victim\_page* (i.e.,  $\delta$  in Algorithm 1) specifies the target number of victim pages that should be maintained in the new victim page list.
- (2) Preform timely page shrinking.

### Algorithm 1: Always Ready Page Shrinking

```
Input: nr_victim_page: the actual number of pages in the victim page list (The initial value is 0);  

nr_ideal_victim_page: the ideal number of pages in the victim page list (The initial value is  $\delta$ );  

shrink_size: the number of page shrinking each time ;  

if System Start then  

  while nr_victim_page < nr_ideal_victim_page do  

    kshrinkd starts;  

    shrink_size = nr_ideal_victim_page ;  

    kshrinkd sleep;
```

Once page writeback is required, provide victim pages;  
**if** Memory Swapping or Direct Reclaim Start **then**  
**while** *nr\_victim\_page* < *nr\_ideal\_victim\_page* **do**  
*kshrinkd* starts;  
*shrink\_size* = *nr\_to\_reclaim* ;





# PMR's Design

## Storage-friendly Page Writeback(SPW)

**Application-aware Page Unmap:** It performs unified page unmap based on applications to quickly serve write I/Os.

**Batch Write I/Os:** It adopts most cost-effective I/O size based on the characteristics of flash-based storage devices.



# Evaluation

## Experiment Setup

**Evaluation Platforms and Workloads:** We adopt different real mobile devices and installed the pre-selected 36 applications.

**Evaluated Schemes:** Five schemes are implemented and measured to show the effectiveness of PMR.

Table 1: Mobile Devices Under Evaluation.

Model	Google Pixel 5	Redmi Note 11	Google Pixel 6 pro
CPU	Qualcomm Snapdragon 765G	MediaTek Dimensity 810	Google Tensor
Memory	8 GB	6 GB	12 GB
Storage	128 GB/UFS 2.1	128 GB/UFS 2.2	256 GB/UFS 3.1
System	Android 13 (Kernel 5.10)	Android 13 (Kernel 5.10)	Android 13 (Kernel 5.10)
Announced	2020, September	2022, January	2021, October

Table 2: Applications and automated user interaction.

Category	Foreground Applications	Auto user inputs
Browser	Chrome	Browse/Read posts
Social Network	Facebook, Twitter	Browse/Read posts
Multimedia	YouTube, Tiktok	Watch videos
Business Utility	Amazon, Gmap, Uber, Spotify	Browse and search Listen music
Game	Angry Birds	Play a stage


\*Background applications: Browser (Firefox, Opera), Social Network (WhatsApp, Instagram, Skype, WeChat, LinkedIn), Multimedia (Spotify, MXPlayer, Netflix, Capcut), Online shopping (Taobao, eBay, AliPay, BOA, Paypal), Business Utility (Booking, Gmail, New York Times, BBC News, OfficeMobile, GoogleDrive), and Game (Hill Climb Racing, Boom Beach, ClashRoyale, Call of Duty).

# Evaluation

## Application Response Evaluation

**Significant performance improvement:** The application response time decreases 43.6% compared with OriginalMR.

**Good compatibility:** When PMR and Fleet are used together, application response time is reduced by 38.9% compared to Fleet alone.




# Evaluation


## Memory Reclaim Evaluation

**Reclaim Throughput Improvement:** The peak throughput of PMR is increased by 82.8% and 75.5% respectively compared with OriginalMR and Acclaim.

**Storage Device Utilization:** With PMR, the average throughput on the Pixel 6 pro is 65% of that on the Pixel 5.



Memory reclaim throughput among different schemes on Google Pixel 6 pro





# Evaluation

## Sensitivity Study

**The ideal pages in the victim page list:**  $\delta$  aims to control the capacity of the victim page list, which directly affects the performance of memory reclaim.

**The size of the application-aware page unmap:** The size affects both the performance of the unmap and the efficiency of write I/O.



# Conclusion

**Motivation:** Sluggish memory reclaim is mainly due to the sub-optimal memory reclaim path, including three aspects: (1) Sequential Execution of Page Writeback and Page Shrinking; (2) Inefficient Page Shrinking; (3) Internally Blocked Page Writeback.

**Design:** PMR parallelizes key steps of memory reclaim, including Proactive Page Shrinking and Storage-friendly Page Writeback.

**Evaluation:** Evaluations on real mobile devices with popular application show PMR achieves significant performance improvement.



# THANKS