

Uncached buffer IO 探索与 在 f2fs 上的支持

韩棋

工程师

vivo 存储系统工



目录

CONTENTS

1

Uncached buffer I/O 背景简介

2

Uncached buffer I/O 技术原理

3

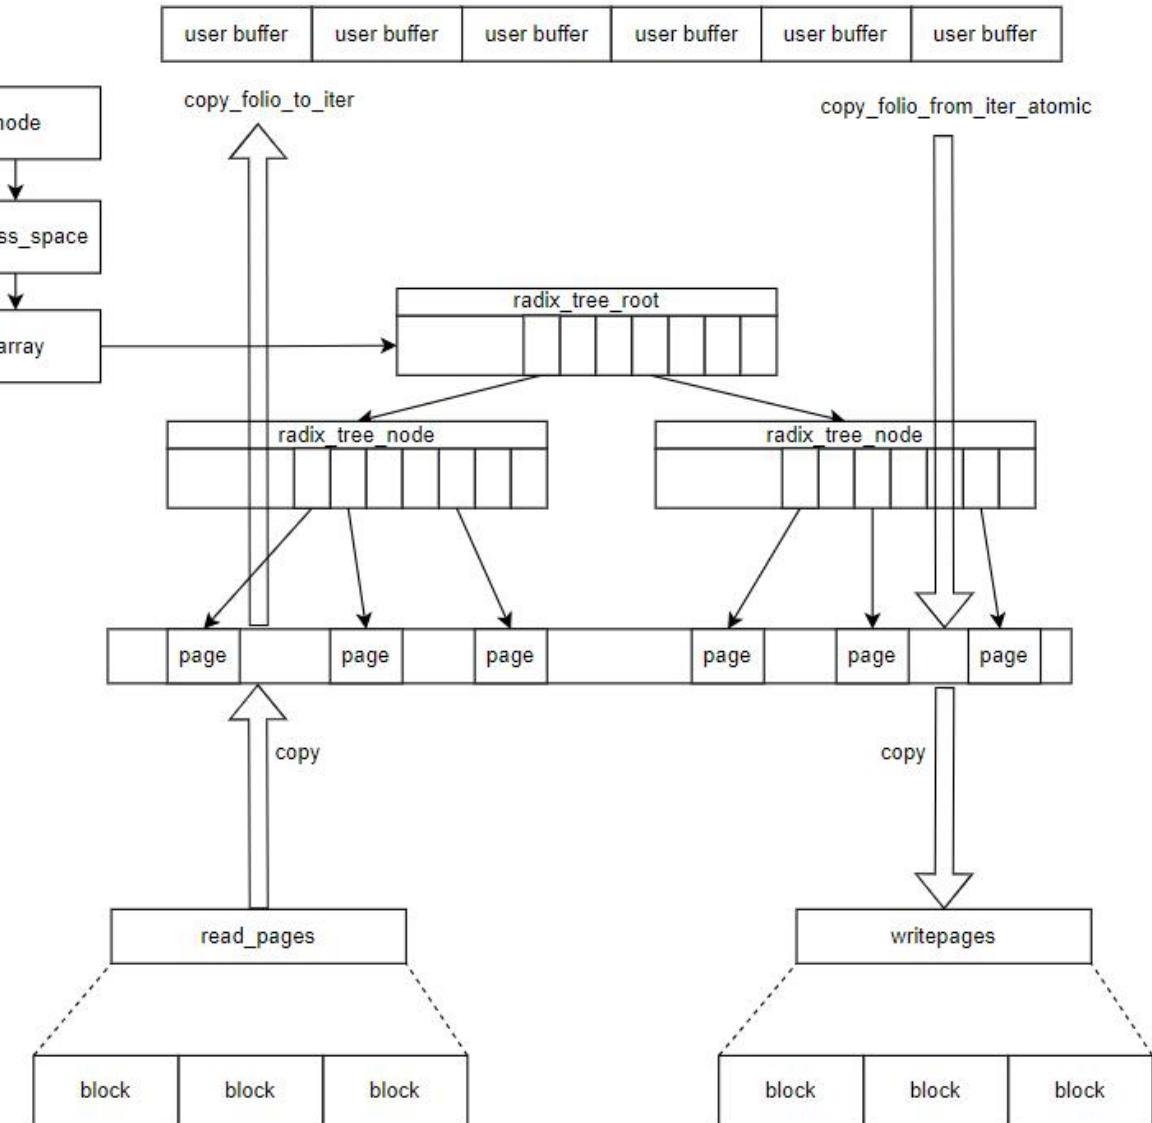
F2FS 适配 Uncached buffer I/O

4

Uncached buffer I/O 收益

1、buffer I/O 可能会消耗大量 CPU 资源进而影响系统性能和功耗

在传统的 buffer I/O 模式中，数据会被存储在内核的页缓存中，以加速未来的访问。然而，当读写大量数据时，页缓存可能会迅速填满，导致内存回收线程（如 kswapd）频繁运行，从而消耗大量 CPU 资源，影响系统性能和功耗。



buffer I/O 读场景下 kswapd 回收线程负载高，消耗 CPU 资源

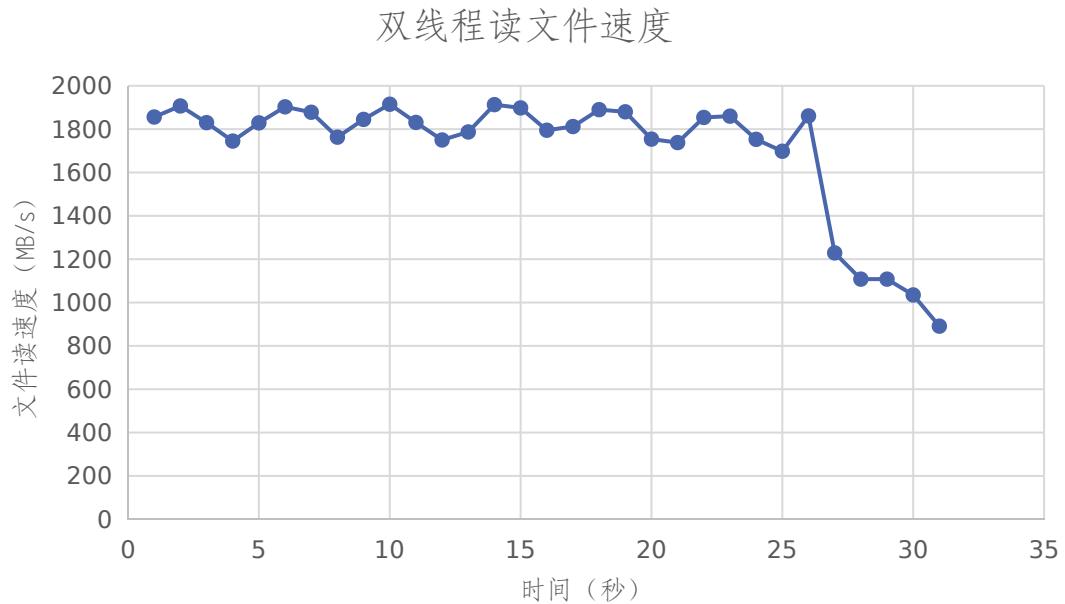


Linux 6.12.22-android16-3-gaaf1b8f85696-4k (localhost) 01/02/25 _aarch64_								Linux 6.12.22-android16-3-gaaf1b8f85696-4k (localhost) 01/02/25 _aarch64_ (8 CPU)											
08:36:26	UID	PID	%usr	%system	%guest	%wait	%CPU	CPU	Command	08:36:24	ppgpin/s	ppgout/s	fault/s	majflt/s	pgfree/s	pgscank/s	pgscand/s	pgsteal/s	ppg
08:36:27	0	93	0.00	0.00	0.00	0.00	0.00	7	kswapd0	08:36:25	1473.27	0.00	99.01	41.58	79.21	0.00	0.00	0.00	0.00
08:36:28	0	93	0.00	0.00	0.00	0.00	0.00	7	kswapd0	08:36:26	0.00	0.00	8.00	1.00	37.00	0.00	0.00	0.00	0.00
08:36:29	0	93	0.00	0.00	0.00	0.00	0.00	7	kswapd0	08:36:27	244.00	0.00	544.00	57.00	375.00	0.00	0.00	0.00	0.00
08:36:30	0	93	0.00	56.00	0.00	0.00	56.00	7	kswapd0	08:36:28	0.00	0.00	0.99	0.00	43.56	0.00	0.00	0.00	0.00
08:36:31	0	93	0.00	73.00	0.00	0.00	73.00	7	kswapd0	08:36:29	4.00	0.00	29.00	0.00	13.00	0.00	0.00	0.00	0.00
08:36:32	0	93	0.00	83.00	0.00	0.00	83.00	7	kswapd0	08:36:30	195968.00	36.00	506.00	0.00	37575.00	37682.00	0.00	74820.00	0.00
08:36:33	0	93	0.00	75.00	0.00	0.00	75.00	7	kswapd0	08:36:31	1884664.00	6932.00	15.00	6.00	477229.00	479930.00	0.00	953362.00	0.00
08:36:34	0	93	0.00	81.00	0.00	0.00	81.00	7	kswapd0	08:36:32	1907200.00	1432.00	7.00	0.00	469891.00	473621.00	0.00	939692.00	0.00
08:36:35	0	93	0.00	54.00	0.00	1.00	54.00	2	kswapd0	08:36:33	1795332.00	1352.00	91.00	1.00	466257.00	484989.00	0.00	932416.00	0.00
08:36:36	0	93	0.00	61.00	0.00	0.00	61.00	0	kswapd0	08:36:34	1734152.00	3100.00	16.00	2.00	431761.00	434012.00	0.00	863498.00	0.00
08:36:37	0	93	0.00	68.00	0.00	0.00	68.00	7	kswapd0	08:36:35	1874692.00	3128.00	5.00	1.00	457221.00	460185.00	0.00	914382.00	0.00
08:36:38	0	93	0.00	53.00	0.00	0.00	53.00	2	kswapd0	08:36:36	1896192.00	3384.00	0.00	0.00	492114.00	494343.00	0.00	984194.00	0.00
08:36:39	0	93	0.00	82.00	0.00	0.00	82.00	7	kswapd0	08:36:37	1866796.00	2668.00	474.00	12.00	454563.00	456222.00	0.00	908322.00	0.00
08:36:40	0	93	0.00	77.00	0.00	0.00	77.00	1	kswapd0	08:36:38	1763412.00	2596.00	35.00	14.00	446324.00	463630.00	0.00	892402.00	0.00
08:36:41	0	93	0.00	74.00	0.00	1.00	74.00	7	kswapd0	08:36:39	1862684.00	1604.00	22.00	0.00	450718.00	452823.00	0.00	901408.00	0.00
08:36:42	0	93	0.00	71.00	0.00	0.00	71.00	7	kswapd0	08:36:40	1912636.00	1012.00	59.00	7.00	483042.00	485815.00	0.00	965788.00	0.00
08:36:43	0	93	0.00	78.00	0.00	0.00	78.00	7	kswapd0	08:36:41	1805068.00	252.00	781.00	0.00	459968.00	462066.00	0.00	918324.00	0.00
08:36:44	0	93	0.00	85.00	0.00	0.00	85.00	7	kswapd0	08:36:42	1753192.00	292.00	204.00	14.00	444417.00	462042.00	0.00	888710.00	0.00
08:36:45	0	93	0.00	83.00	0.00	0.00	83.00	7	kswapd0	08:36:43	1806640.00	868.00	102.00	3.00	437412.00	440025.00	0.00	874778.00	0.00
08:36:46	0	93	0.00	70.00	0.00	0.00	70.00	7	kswapd0	08:36:44	1915996.00	1188.00	14.00	3.00	485006.00	488247.00	0.00	969954.00	0.00
08:36:47	0	93	0.00	78.00	0.00	1.00	78.00	2	kswapd0	08:36:45	1888060.00	636.00	6.00	2.00	482974.00	485947.00	0.00	965790.00	0.00
08:36:48	0	93	0.00	81.00	0.00	0.00	81.00	3	kswapd0	08:36:46	1794816.00	812.00	50.00	0.00	441887.00	459891.00	0.00	883678.00	0.00
08:36:49	0	93	0.00	54.00	0.00	0.00	54.00	7	kswapd0	08:36:47	1822480.00	1396.00	34.00	9.00	452257.00	454652.00	0.00	904432.00	0.00
08:36:50	0	93	0.00	76.00	0.00	0.00	76.00	1	kswapd0	08:36:48	1898056.00	956.00	112.00	11.00	480725.00	483584.00	0.00	961344.00	0.00
08:36:51	0	93	0.00	75.00	0.00	0.00	75.00	0	kswapd0	08:36:49	1841264.00	1396.00	24.00	1.00	457446.00	459797.00	0.00	914844.00	0.00
08:36:52	0	93	0.00	73.00	0.00	0.00	73.00	7	kswapd0	08:36:50	1777428.00	2316.00	466.00	143.00	447691.00	464679.00	0.00	894362.00	0.00
08:36:53	0	93	0.00	61.00	0.00	1.00	61.00	7	kswapd0	08:36:51	1752588.00	3292.00	2.00	0.00	447774.00	449608.00	0.00	895476.00	0.00
08:36:54	0	93	0.00	80.00	0.00	0.00	80.00	7	kswapd0	08:36:52	1852696.00	1052.00	4.00	1.00	460700.00	463014.00	0.00	921216.00	0.00
08:36:55	0	93	0.00	64.00	0.00	0.00	64.00	7	kswapd0	08:36:53	1829572.00	1044.00	107.00	6.00	441776.00	444011.00	0.00	883488.00	0.00
08:36:56	0	93	0.00	56.00	0.00	0.00	56.00	7	kswapd0	08:36:54	1770812.00	880.00	7.00	2.00	449181.00	466480.00	0.00	898270.00	0.00
08:36:57	0	93	0.00	26.00	0.00	0.00	26.00	2	kswapd0	08:36:55	1714484.00	1844.00	7.00	5.00	425720.00	427673.00	0.00	851370.00	0.00
08:36:58	0	93	0.00	24.00	0.00	1.00	24.00	3	kswapd0	08:36:56	1840228.00	2364.00	5.00	0.00	469578.00	471321.00	0.00	939122.00	0.00
08:36:59	0	93	0.00	22.00	0.00	1.00	22.00	3	kswapd0	08:36:57	1186616.00	428.00	52.00	12.00	293856.00	294889.00	0.00	587688.00	0.00
08:37:00	0	93	0.00	15.84	0.00	0.00	15.84	3	kswapd0	08:36:58	1122412.00	368.00	2248.00	1427.00	277517.00	277487.00	0.00	552644.00	0.00
08:37:01	0	93	0.00	0.00	0.00	0.00	0.00	3	kswapd0	08:36:59	1111904.00	524.00	32.00	3.00	287897.00	289377.00	0.00	575702.00	0.00
08:37:02	0	93	0.00	0.00	0.00	0.00	0.00	3	kswapd0	08:37:00	1021036.00	484.00	134.00	22.00	260484.00	276913.00	0.00	520398.00	0.00
										08:37:01	278200.00	256.00	36.00	4.00	68447.00	68659.00	0.00	136492.00	0.00
										08:37:02	160.00	0.00	209.00	40.00	34.00	0.00	0.00	0.00	0.00
										Average:	1380304.37	1312.26	172.23	48.68	345607.34	350173.96	0.00	690909.15	0.00

buffer I/O 读文件速度不稳定，本地环境存在 150MB/s 的速度波动

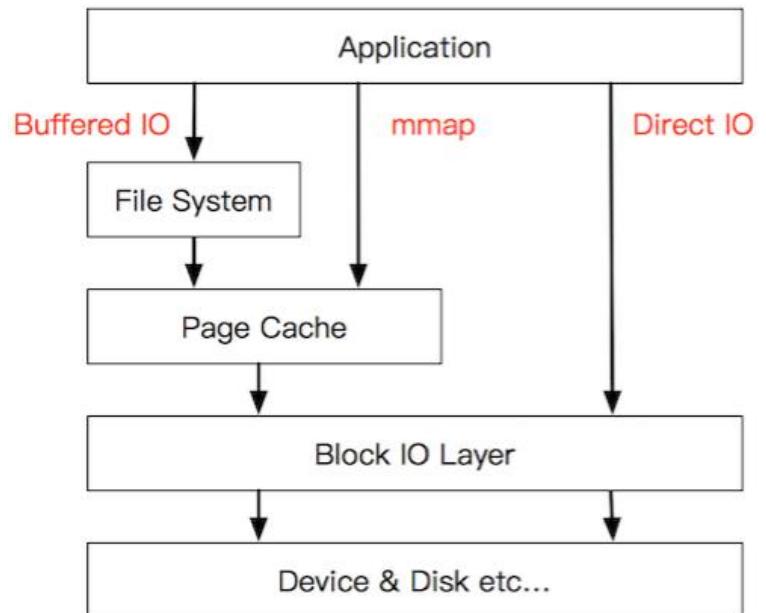
普通读

```
cmd: uncached_io_test 32768 0 1 25G.txt 25_1G.txt
Starting 2 threads
reading bs 32768, uncached 0
 1s: 1856MB/sec, MB=1856
 2s: 1907MB/sec, MB=3763
 3s: 1830MB/sec, MB=5594
 4s: 1745MB/sec, MB=7333
 5s: 1829MB/sec, MB=9162
 6s: 1903MB/sec, MB=11075
 7s: 1878MB/sec, MB=12942
 8s: 1763MB/sec, MB=14718
 9s: 1845MB/sec, MB=16549
10s: 1915MB/sec, MB=18481
11s: 1831MB/sec, MB=20295
12s: 1750MB/sec, MB=22066
13s: 1787MB/sec, MB=23832
14s: 1913MB/sec, MB=25769
15s: 1898MB/sec, MB=27668
16s: 1795MB/sec, MB=29436
17s: 1812MB/sec, MB=31248
18s: 1890MB/sec, MB=33139
19s: 1880MB/sec, MB=35020
20s: 1754MB/sec, MB=36810
21s: 1738MB/sec, MB=38511
22s: 1854MB/sec, MB=40404
23s: 1860MB/sec, MB=42224
24s: 1753MB/sec, MB=43978
25s: 1698MB/sec, MB=45676
26s: 1861MB/sec, MB=47584
27s: 1229MB/sec, MB=48766
28s: 1108MB/sec, MB=49923
29s: 1108MB/sec, MB=51031
30s: 1035MB/sec, MB=52016
30s: 891MB/sec, MB=129774
```



direct I/O 使用有严格内存对齐和大小限制

Direct I/O 虽然可以绕过 page cache，直接在用户空间和存储设备之间传输数据，进而避免了内存消耗，但这种方式要求严格的内存对齐和大小限制，使用起来较为复杂，且在 Android 上层无法直接被使用。



O_DIRECT

The O_DIRECT flag may impose alignment restrictions on the length and address of user-space buffers and the file offset of I/Os. In Linux alignment restrictions vary by filesystem and kernel version and might be absent entirely. However there is currently no filesystem-independent interface for an application to discover these restrictions for a given file or filesystem. Some filesystems provide their own interfaces for doing so, for example the XFS_IOC_DIOINFO operation in xfsctl(3).

Under Linux 2.4, transfer sizes, and the alignment of the user buffer and the file offset must all be multiples of the logical block size of the filesystem. Since Linux 2.6.0, alignment to the logical block size of the underlying storage (typically 512 bytes) suffices. The logical block size can be determined using the ioctl(2) BLKSSZGET operation or from the shell using the command:

```
blockdev --getss
```

2、Uncached buffer I/O 诞生

为了能够使用 Buffer I/O 的便利性，同时又让 Buffer I/O 不占用 page cache，Jens Axboe 提出了 Uncached buffered I/O。

The screenshot shows a portion of the LWN.net website. At the top right, there are user login fields ('User:' and 'Password:'), a 'Log in' button, and links for 'Subscribe' and 'Register'. Below this, the title 'Uncached buffered IO' is displayed. The main content is an email message from Jens Axboe. The email's 'From' field is highlighted with a red border. The message body discusses the history of support for RWF_UNCACHED and introduces a new approach using page cache synchronization for buffered I/O.

User: **Password:** **Log in** | **Subscribe** | **Register**

Uncached buffered IO

From: [Jens Axboe <axboe-AT-kernel.dk>](#)
To: [linux-mm-AT-kvack.org](#), [linux-fsdevel-AT-vger.kernel.org](#)
Subject: [PATCHSET v4] Uncached buffered IO
Date: Fri, 08 Nov 2024 10:43:23 -0700
Message-ID: <20241108174505.1214230-1-axboe@kernel.dk>
Cc: [hannes-AT-cmpxchg.org](#), [clm-AT-meta.com](#), [linux-kernel-AT-vger.kernel.org](#)
Archive-link: [Article](#)

Hi,

5 years ago I posted patches adding support for RWF_UNCACHED, as a way to do buffered I/O that isn't page cache persistent. The approach back then was to have private pages for I/O, and then get rid of them once I/O was done. But that then runs into all the issues that O_DIRECT has, in terms of synchronizing with the page cache.

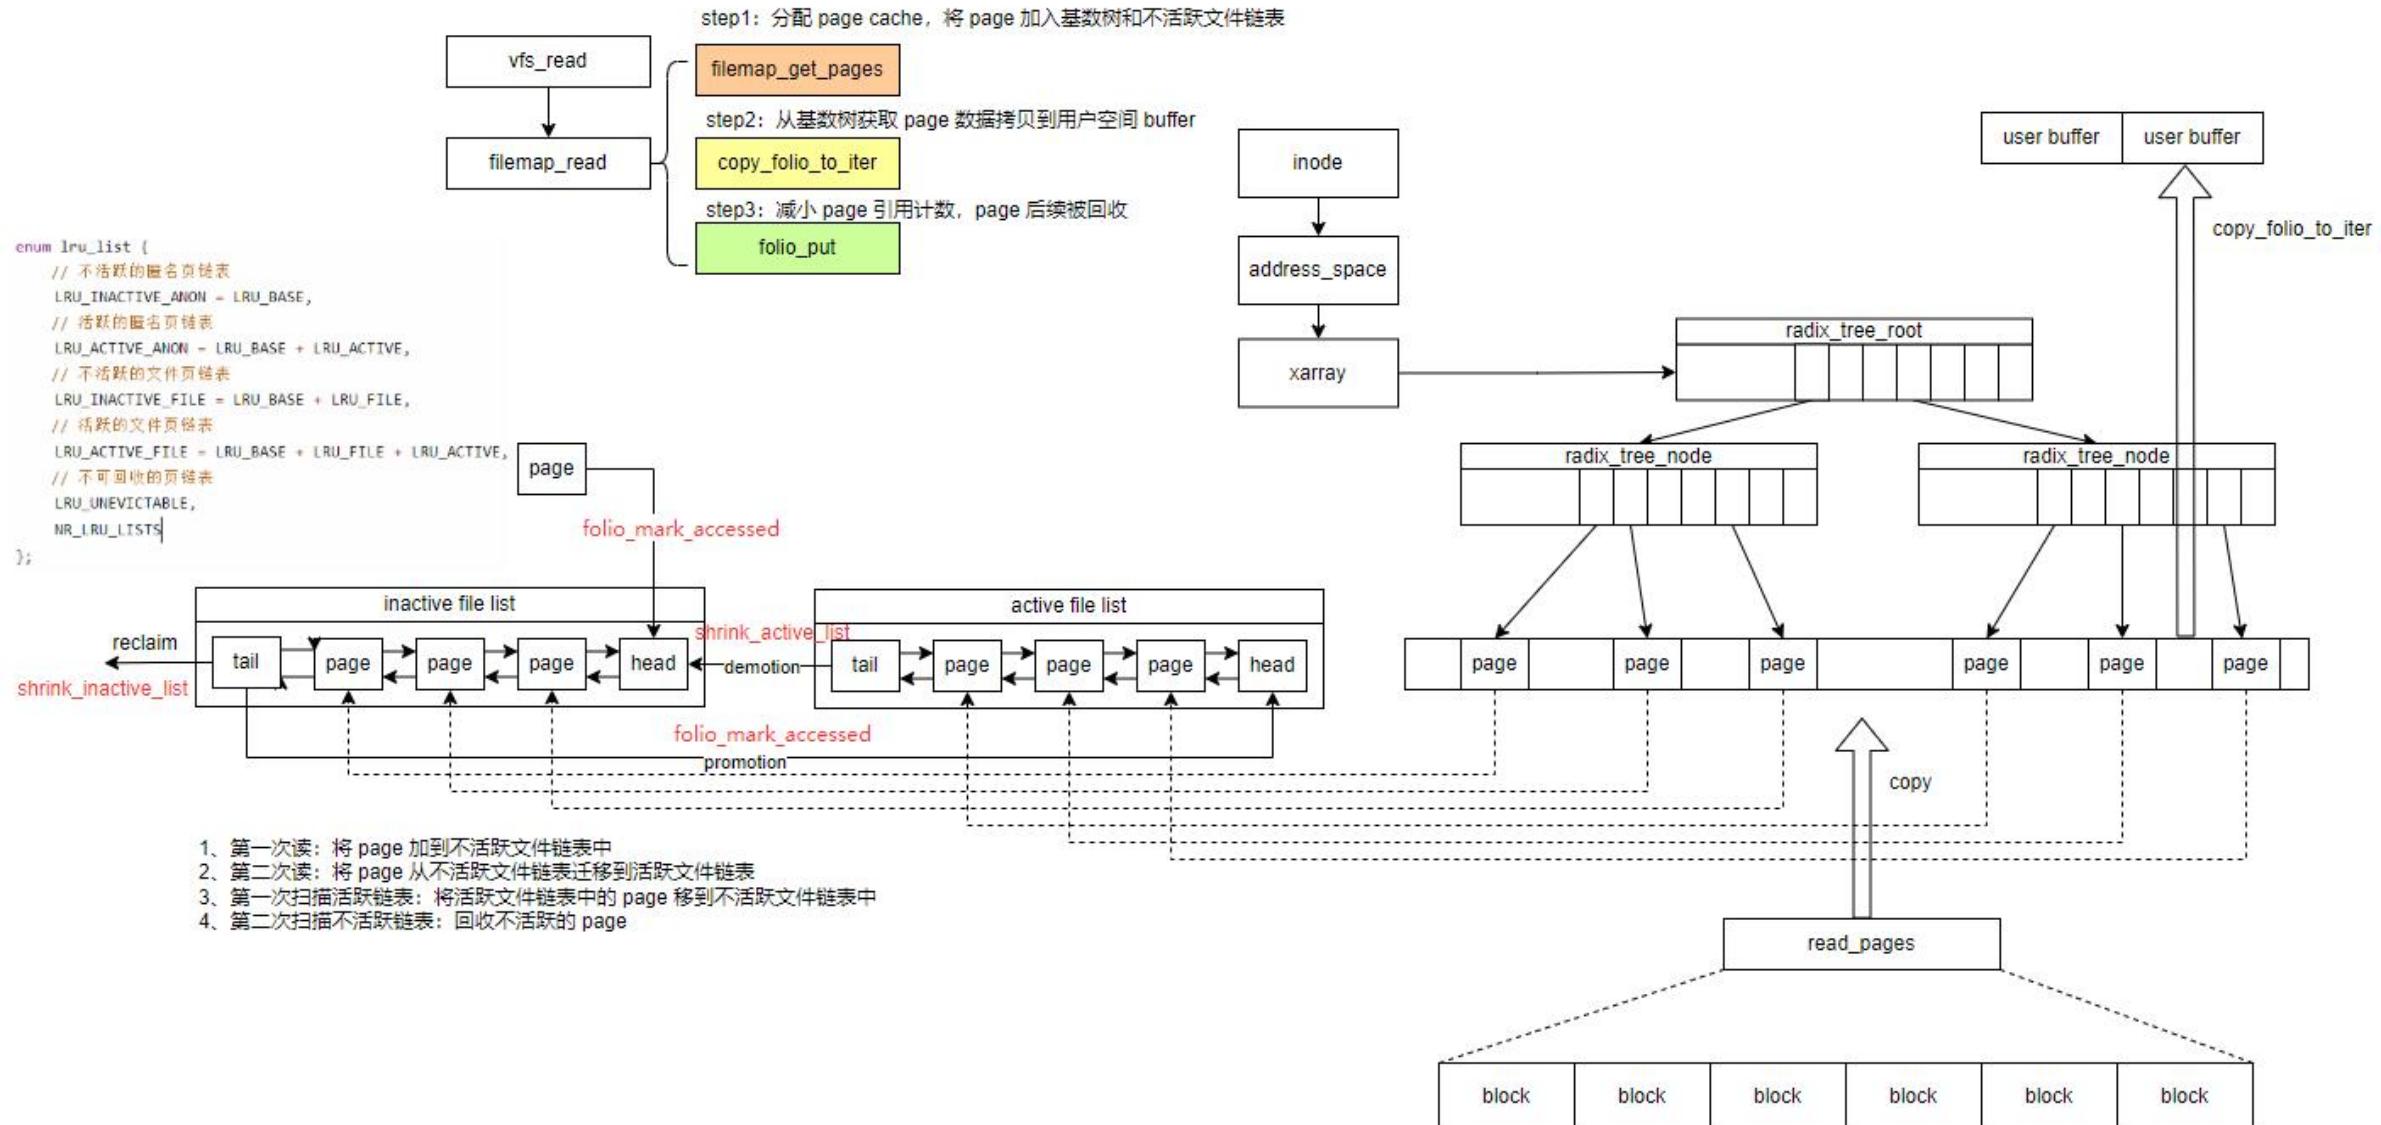
So here's a new approach to the same concern, but using the page cache as synchronization. That makes RWF_UNCACHED less special, in that it's just page cache I/O, except it prunes the ranges once I/O is completed.

Why do this, you may ask? The tl;dr is that device speeds are only getting faster, while reclaim is not. Doing normal buffered I/O can be very unpredictable, and suck up a lot of resources on the reclaim side. This leads people to use O_DIRECT as a work-around, which has its own set of restrictions in terms of size, offset, and length of I/O. It's also inherently synchronous, and now you need async I/O as well. While the latter isn't necessarily a big problem as we have good options available there, it also should not be a requirement when all you want to do is read or write some data without caching.

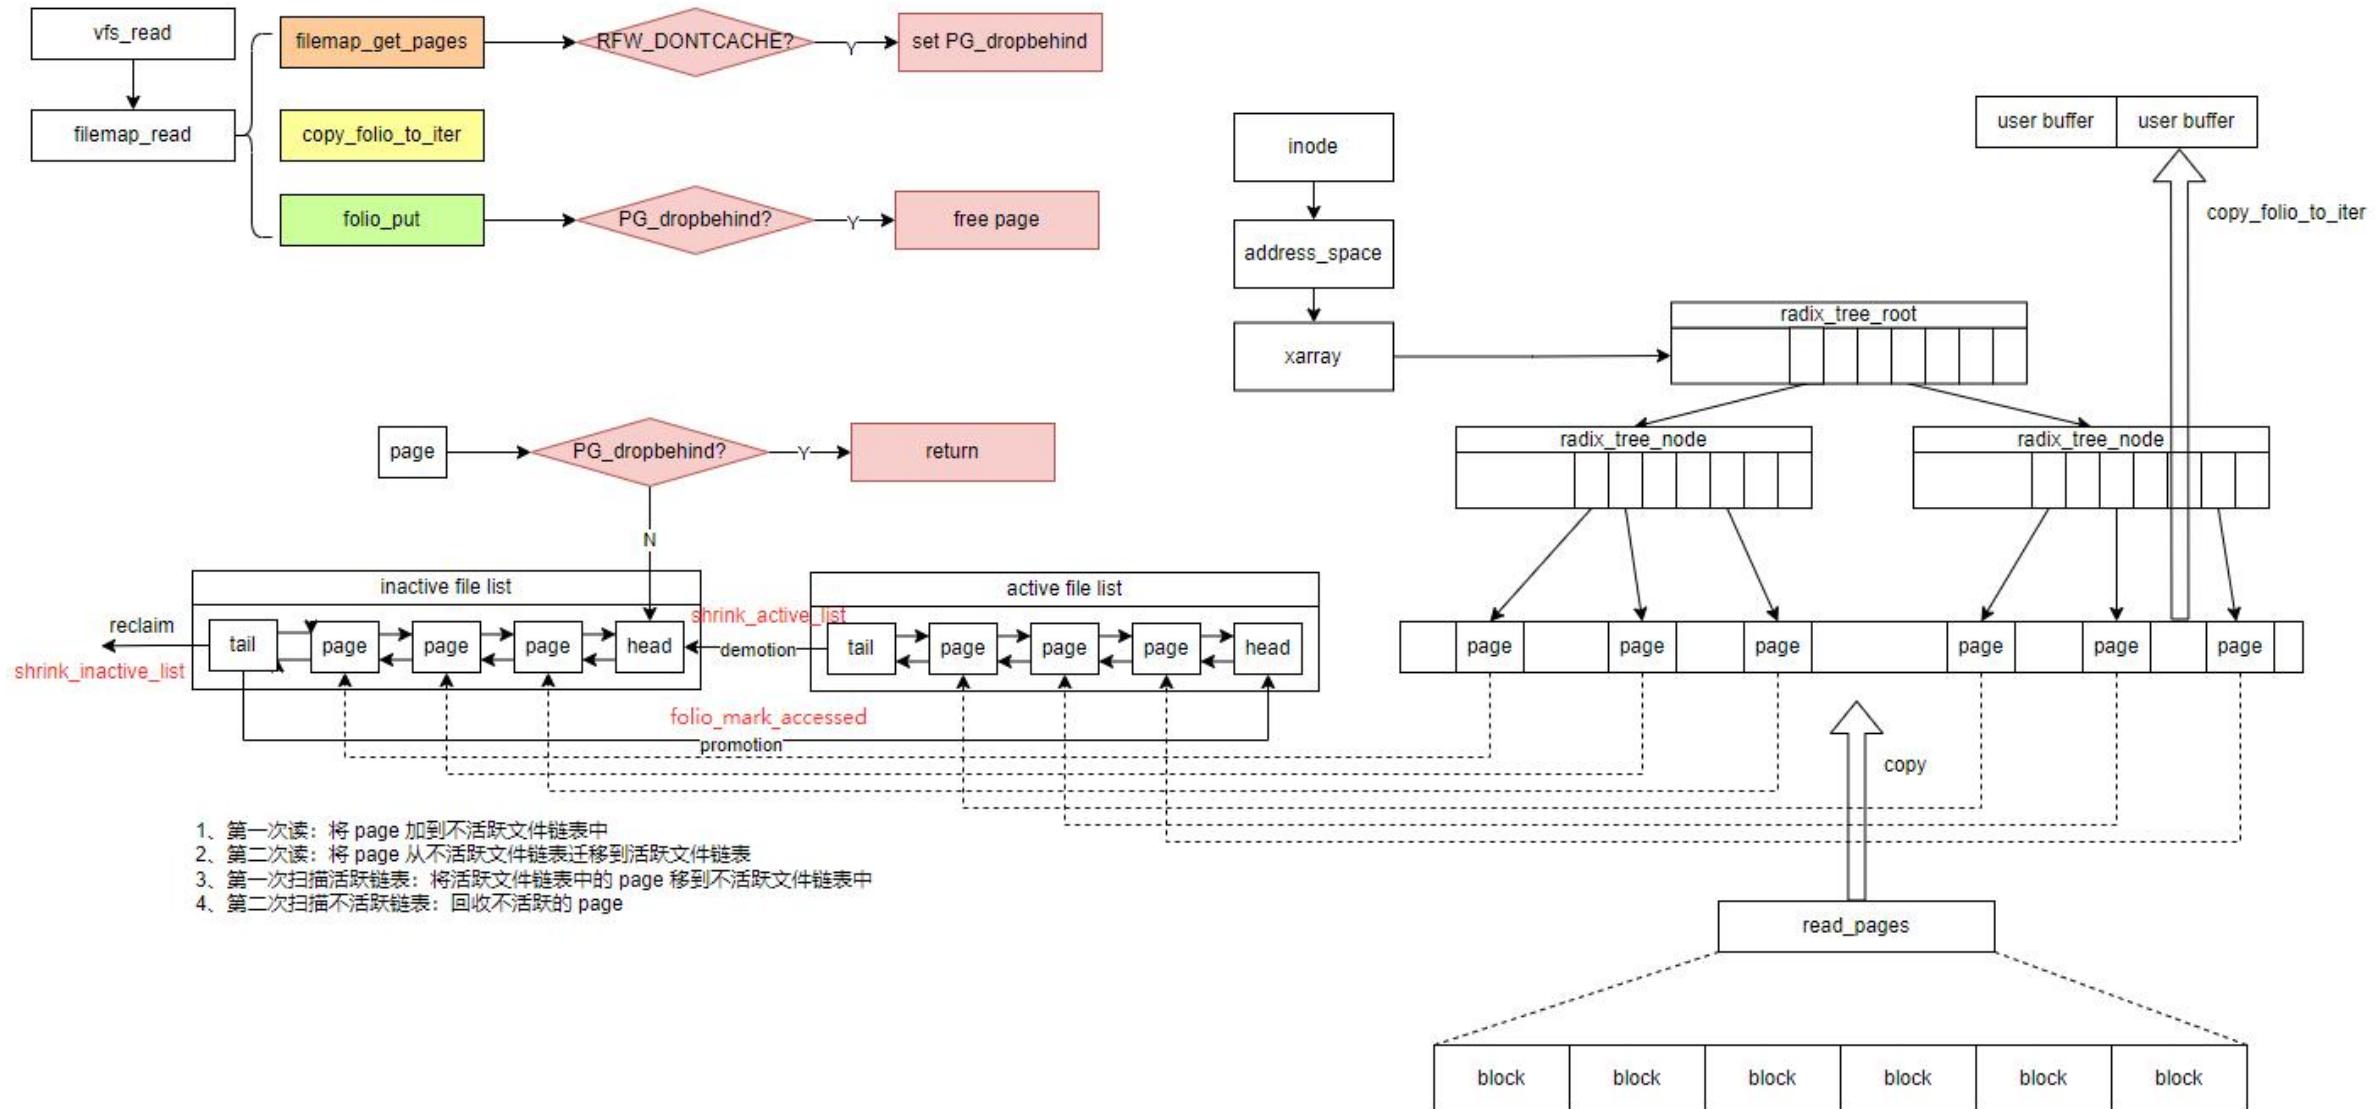
Even on desktop type systems, a normal NVMe device can fill the entire page cache in seconds. On the big system I used for testing, there's a lot more RAM, but also a lot more devices. As can be seen in some of the results in the following patches, you can still fill RAM in seconds even when there's 1TB of it. Hence this problem isn't solely a "big hyperscaler system" issue, it's common across the board. Normal users do big backups too, edit videos, etc.

Common for both reads and writes with RWF_UNCACHED is that they use the page cache for I/O. Reads work just like a normal buffered read would, with the only exception being that the touched ranges will get pruned after data has been copied. For writes, the ranges will get writeback kicked off before the syscall returns, and then writeback completion will prune the range. Hence writes aren't synchronous, and it's easy to pipeline writes using RWF_UNCACHED.

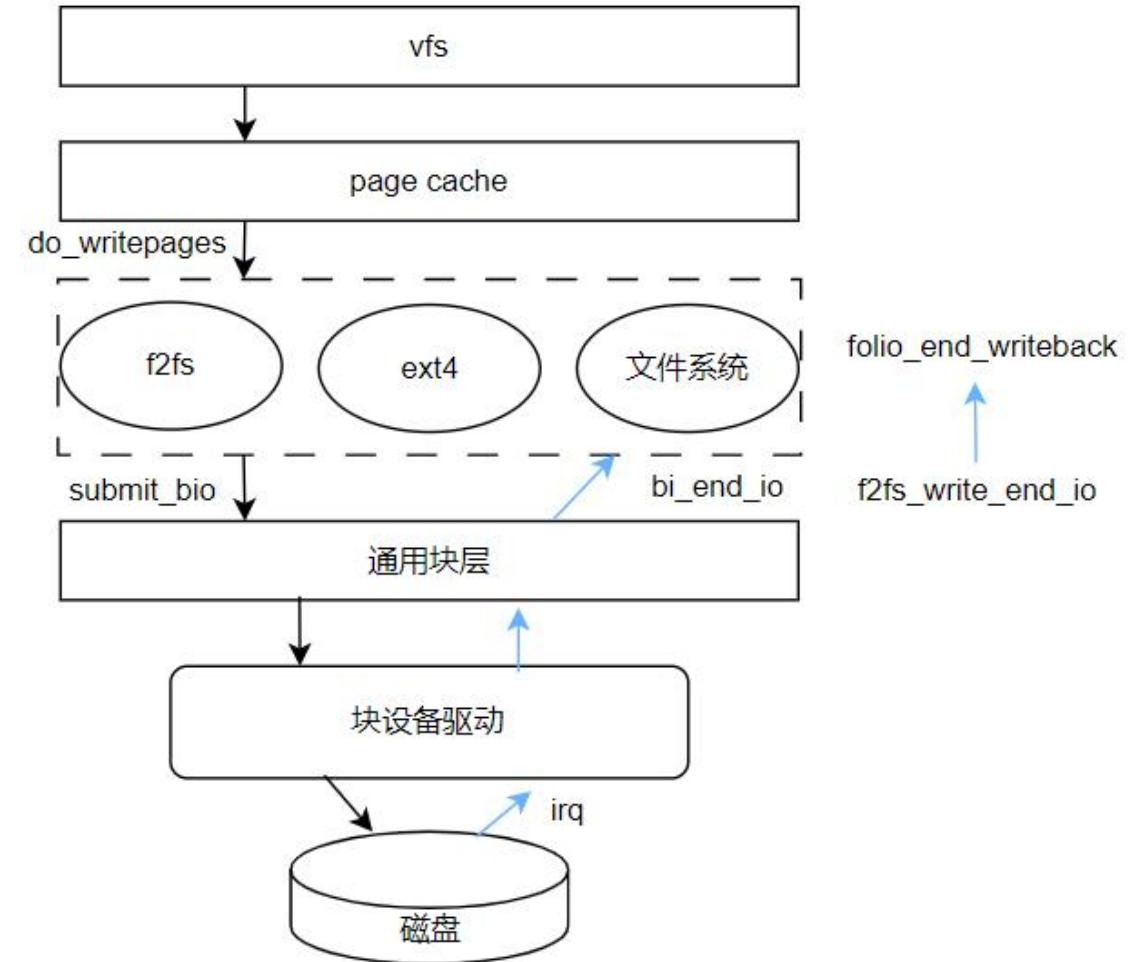
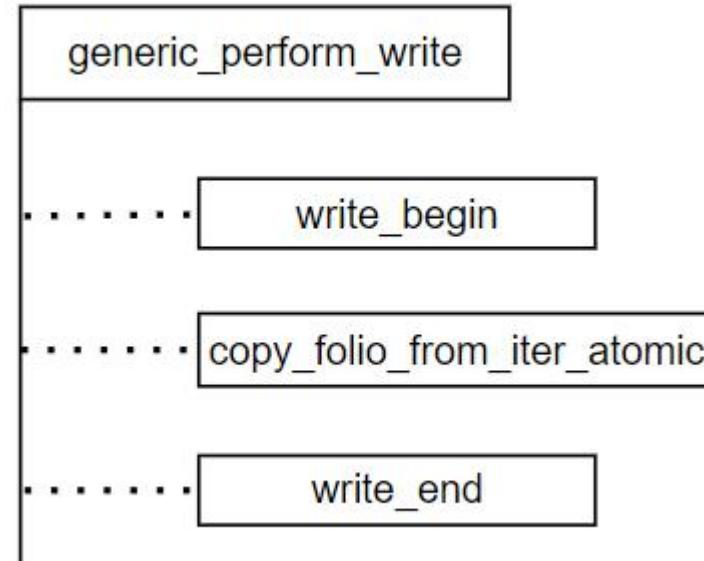
内核通用读文件与 page cache 回收流程



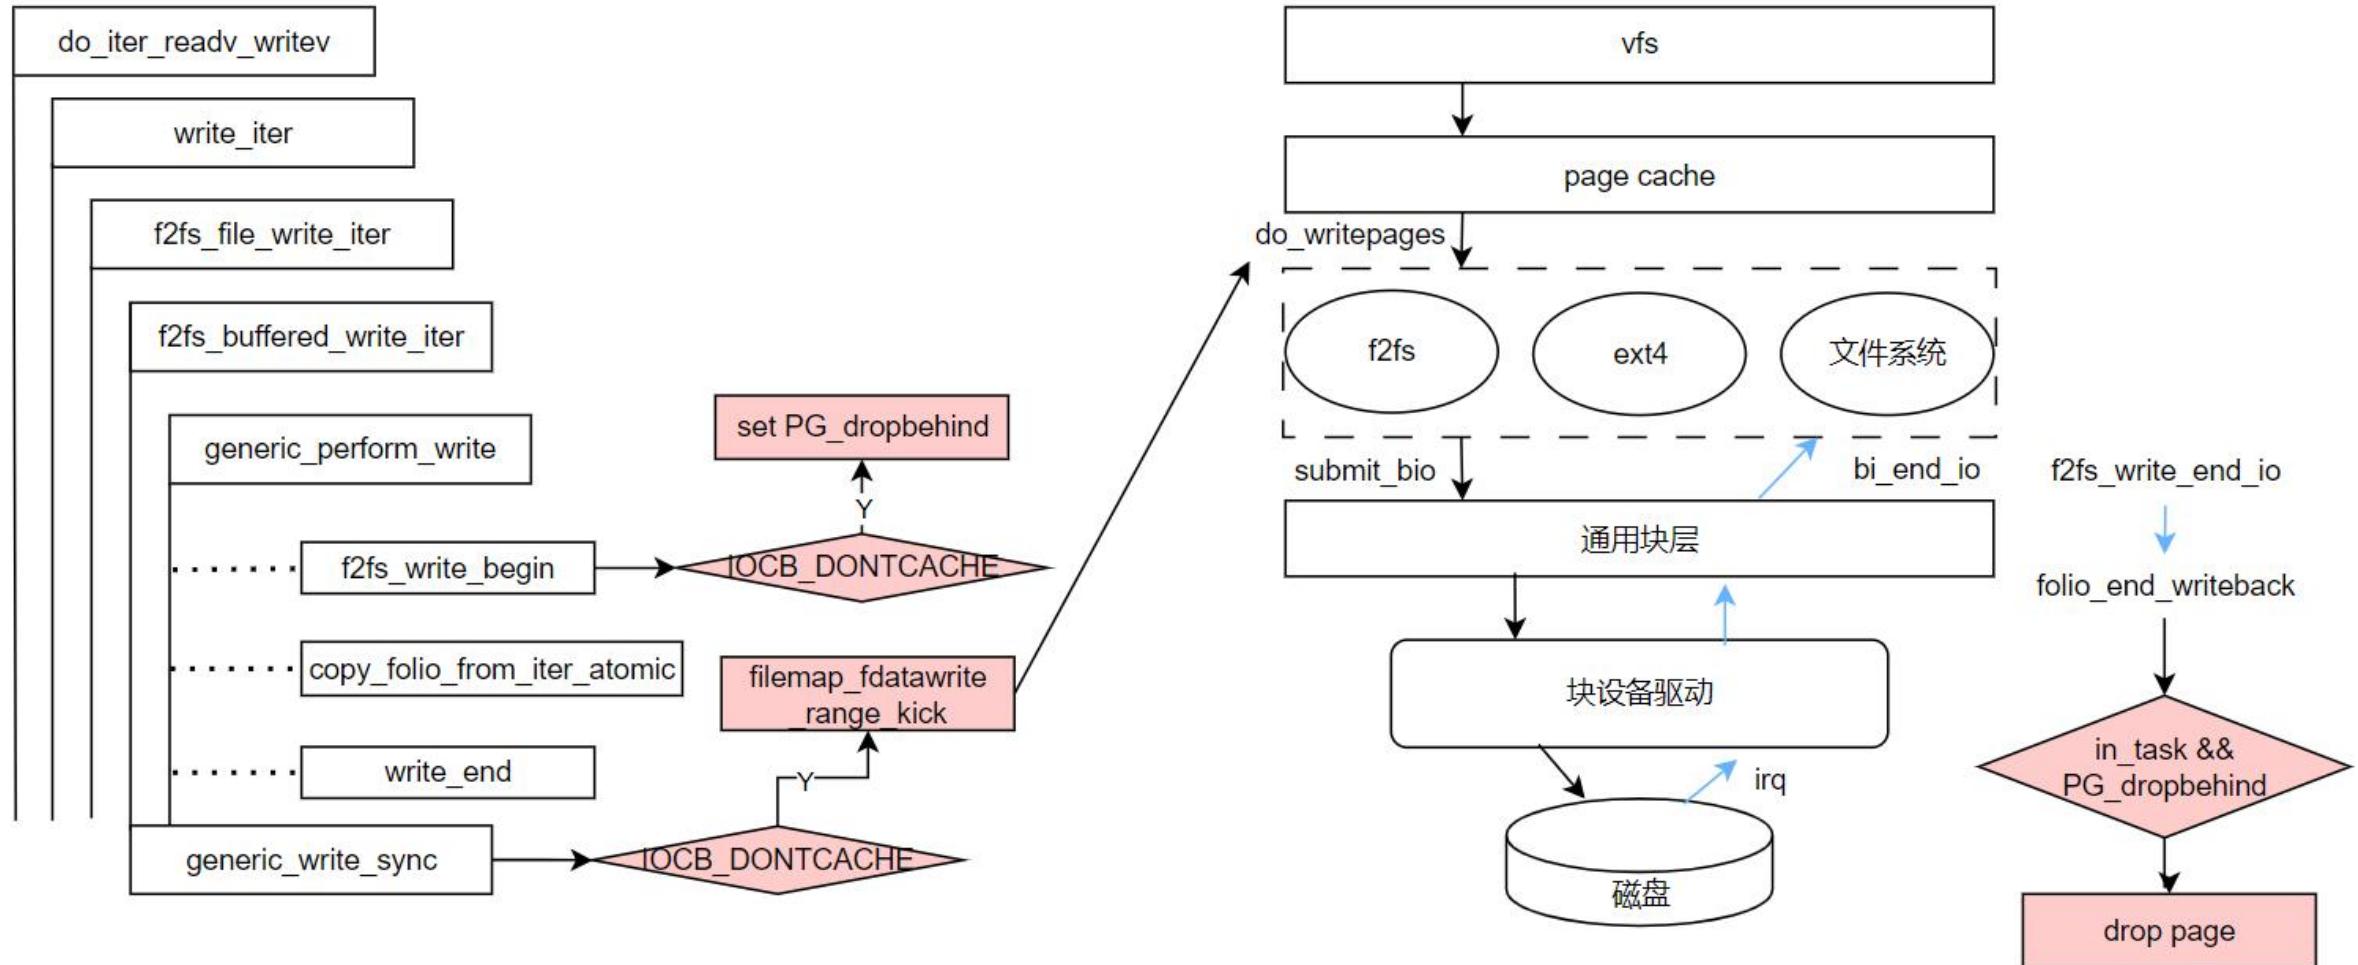
uncached buffer io 读文件实现原理



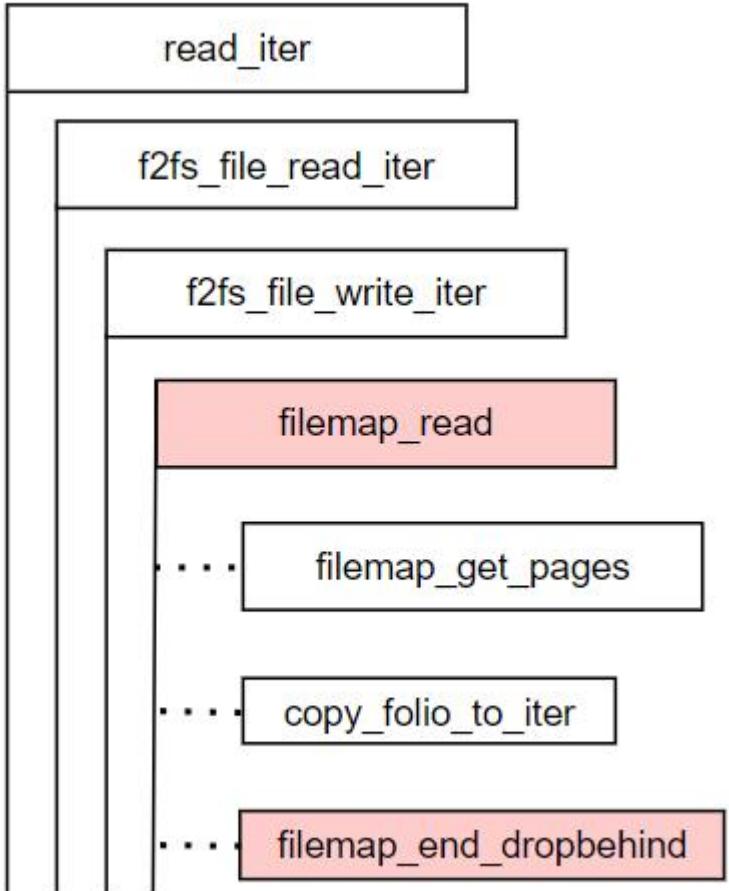
内核通用写文件与 page cache 回收流程



uncached buffer io 写文件实现原理

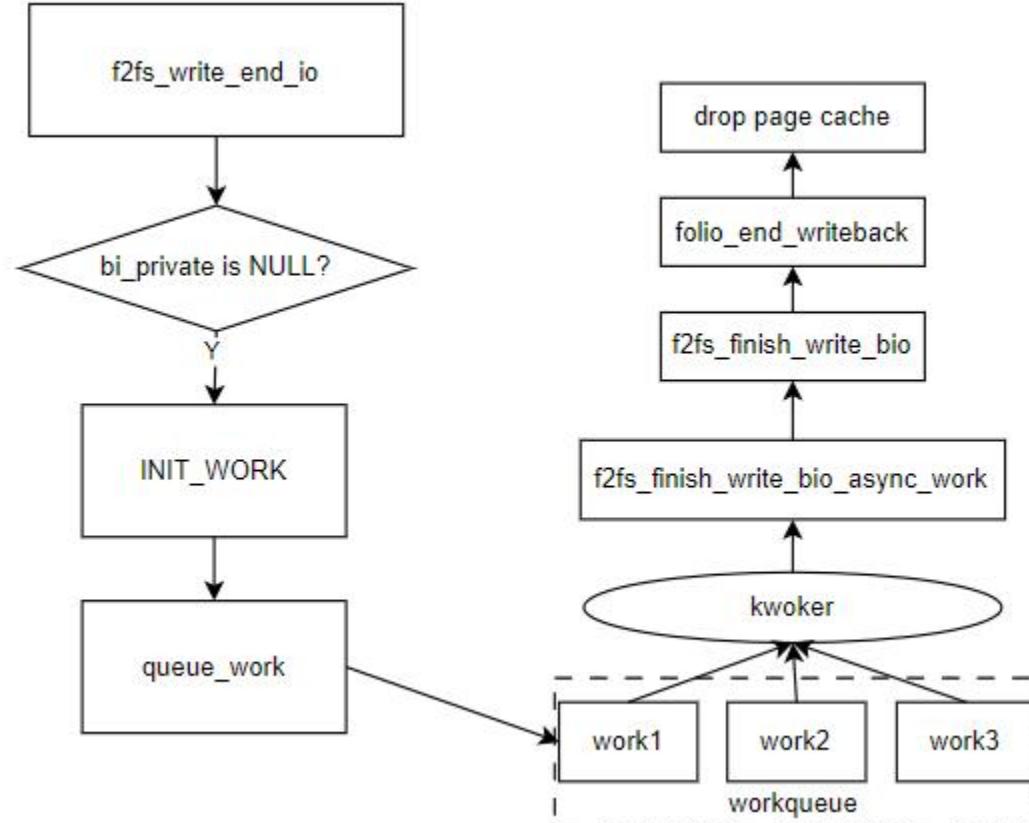
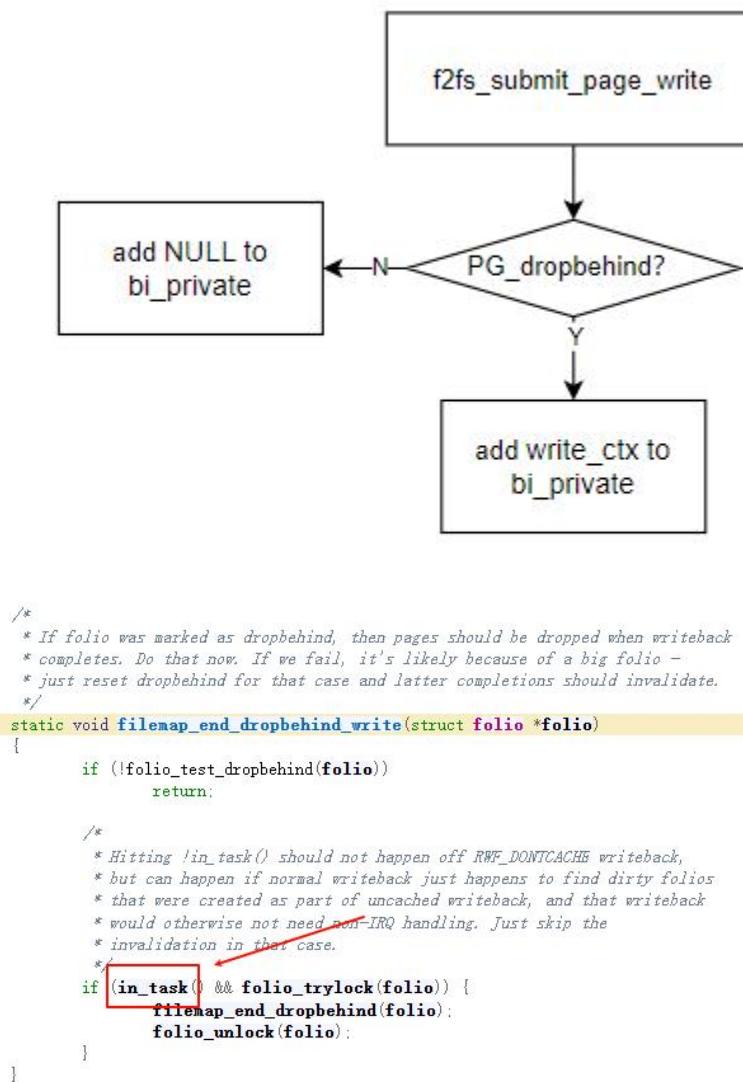


3、f2fs 适配 Uncached buffer I/O 读



f2fs 适配 Uncached buffer IO 读提交链接：
<https://lore.kernel.org/all/20250725075310.1614614-1-hanqi@vivo.com/>

f2fs 适配 Uncached buffer I/O 写



f2fs 适配 Uncached buffer IO 读提交链接：
<https://lore.kernel.org/all/20250828121131.36941>
54-1-hanqi@vivo.com/

4、f2fs Uncached buffer I/O 读场景后台无内存交换， kswapd 负载降为 0

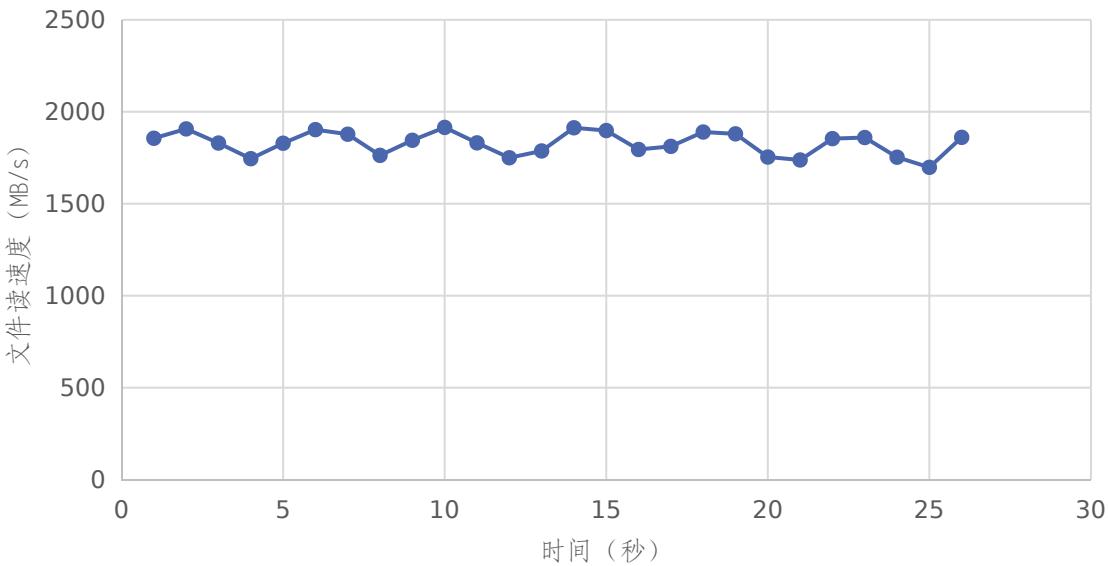


Linux 6.12.22-android16-3-gaaf1b8f85696-4k (localhost) 01/02/25 _aarch64_ (8 CPU)							
08:38:01	pgpgin/s	pgpgout/s	fault/s	majflt/s	pgfree/s	pgscank/s	pgscand/s
08:38:02	1132.00	0.00	179.00	33.00	64.00	0.00	0.00
08:38:03	137896.00	0.00	467.00	10.00	34410.00	0.00	0.00
08:38:04	1876228.00	16.00	97.00	1.00	469070.00	0.00	0.00
08:38:05	1911552.00	0.00	1.00	0.00	477989.00	0.00	0.00
08:38:06	1868468.00	19968.00	3115.00	835.00	467201.00	716.00	70.00
08:38:07	1877552.00	696.00	797.00	103.00	465825.00	0.00	0.00
08:38:08	1865444.00	4.00	69.00	36.00	466506.00	0.00	0.00
08:38:09	1872000.00	88.00	869.00	648.00	475077.00	11247.00	0.00
08:38:10	1877412.00	28.00	205.00	88.00	470808.00	0.00	0.00
08:38:11	1876940.00	64.00	303.00	110.00	469107.00	0.00	0.00
08:38:12	1841184.00	344.00	850.00	107.00	460102.00	0.00	0.00
08:38:13	1853872.00	364.00	179.00	117.00	466939.00	0.00	0.00
08:38:14	1857932.00	16.00	208.00	99.00	463535.00	0.00	0.00
08:38:15	1812292.00	840.00	2930.00	1474.00	465835.00	11223.00	0.00
08:38:16	1900176.00	0.00	61.00	35.00	475151.00	0.00	0.00
08:38:17	1821112.00	0.00	1349.00	545.00	453042.00	0.00	0.00
08:38:18	1902132.00	192.00	781.00	408.00	474311.00	0.00	0.00
08:38:19	1922112.00	24.00	35.00	12.00	480758.00	0.00	0.00
08:38:20	1858120.00	0.00	128.00	14.00	464872.00	0.00	0.00
08:38:21	1907428.00	164.00	583.00	425.00	476716.00	0.00	0.00
08:38:22	1894488.00	0.00	71.00	67.00	473160.00	0.00	0.00
08:38:23	1905948.00	108.00	23.00	23.00	502150.00	22797.00	0.00
08:38:24	1848016.00	16.00	4978.00	881.00	462187.00	0.00	0.00
08:38:25	1803176.00	0.00	55.00	49.00	450709.00	0.00	0.00
08:38:26	1731656.00	364.00	27317.00	15977.00	441384.00	23123.00	118.00
08:38:27	1141780.00	0.00	23.00	5.00	285451.00	0.00	0.00
08:38:28	992004.00	36.00	2259.00	1802.00	246206.00	0.00	0.00
08:38:29	308204.00	0.00	826.00	119.00	77504.00	0.00	0.00
08:38:30	1620.00	0.00	281.00	196.00	36.00	0.00	0.00
08:38:31	60.00	0.00	67.00	15.00	45.00	0.00	0.00
08:38:32	336.00	0.00	72.00	20.00	9.00	0.00	0.00
^C							

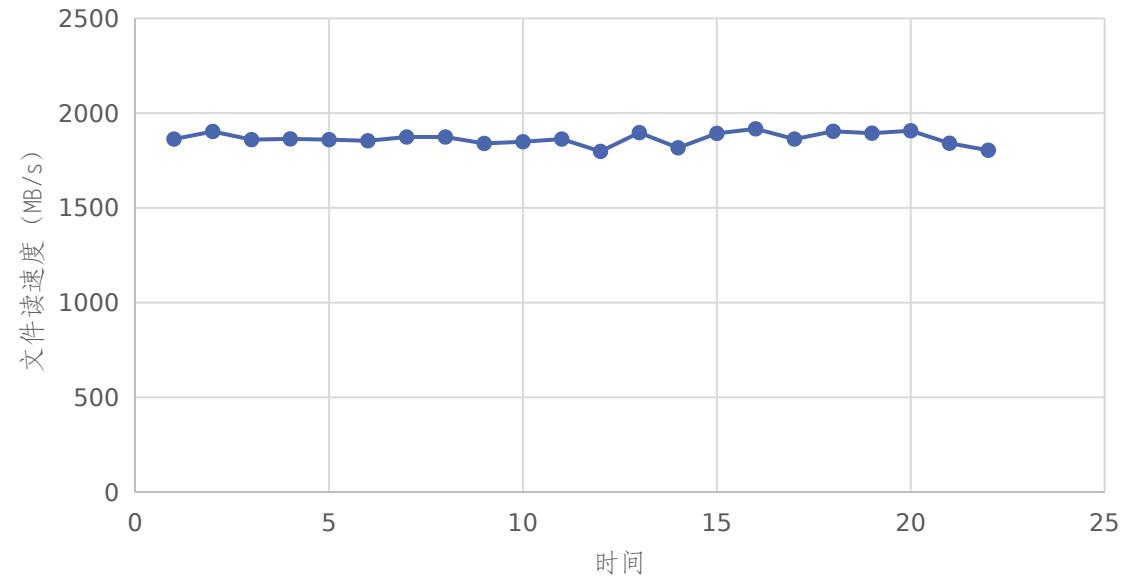
PD2502:/ # pidstat -u -p 93 1							
Linux 6.12.22-android16-3-gaaf1b8f85696-4k (localhost) 01/02/25 _aarch64_							
08:38:00	UID	PID	%usr	%system	%guest	%wait	%CPU
08:38:01	0	93	0.00	0.00	0.00	0.00	0.00
08:38:02	0	93	0.00	0.00	0.00	0.00	0.00
08:38:03	0	93	0.00	0.00	0.00	0.00	0.00
08:38:04	0	93	0.00	0.00	0.00	0.00	0.00
08:38:05	0	93	0.00	0.00	0.00	0.00	0.00
08:38:06	0	93	0.00	1.00	0.00	1.00	1.00
08:38:07	0	93	0.00	0.00	0.00	0.00	0.00
08:38:08	0	93	0.00	0.00	0.00	0.00	0.00
08:38:09	0	93	0.00	1.00	0.00	0.00	1.00
08:38:10	0	93	0.00	0.00	0.00	0.00	0.00
08:38:11	0	93	0.00	0.00	0.00	0.00	0.00
08:38:12	0	93	0.00	0.00	0.00	0.00	0.00
08:38:13	0	93	0.00	0.00	0.00	0.00	0.00
08:38:14	0	93	0.00	0.00	0.00	0.00	0.00
08:38:15	0	93	0.00	3.00	0.00	0.00	3.00
08:38:16	0	93	0.00	0.00	0.00	0.00	0.00
08:38:17	0	93	0.00	0.00	0.00	0.00	0.00
08:38:18	0	93	0.00	0.00	0.00	0.00	0.00
08:38:19	0	93	0.00	0.00	0.00	0.00	0.00
08:38:20	0	93	0.00	0.00	0.00	0.00	0.00
08:38:21	0	93	0.00	0.00	0.00	0.00	0.00
08:38:22	0	93	0.00	0.00	0.00	0.00	0.00
08:38:23	0	93	0.00	3.00	0.00	0.00	3.00
08:38:24	0	93	0.00	0.00	0.00	0.00	0.00
08:38:25	0	93	0.00	0.00	0.00	0.00	0.00
08:38:26	0	93	0.00	4.00	0.00	0.00	4.00
08:38:27	0	93	0.00	0.00	0.00	0.00	0.00
08:38:28	0	93	0.00	0.00	0.00	0.00	0.00
08:38:29	0	93	0.00	0.00	0.00	0.00	0.00
08:38:30	0	93	0.00	0.00	0.00	0.00	0.00
08:38:31	0	93	0.00	0.00	0.00	0.00	0.00
08:38:32	0	93	0.00	0.00	0.00	0.00	0.00
08:38:33	0	93	0.00	0.00	0.00	0.00	0.00

f2fs Uncached buffer I/O 读速度波动由 150MB/s 降到 50MB/s

双线程 buffer IO 读文件速度



双线程 Uncached buffer IO 读文件速度



f2fs Uncached buffer I/O 写场景后台无 page cache 占用，kswapd 负载降为 0



	UID	PID	%usr	%system	%guest	%wait	%CPU	CPU	Command
19:29:34	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:35	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:36	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:37	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:38	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:39	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:40	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:29:41	0	94	0.00	2.00	0.00	0.00	2.00	0	kswapd0
19:29:42	0	94	0.00	59.00	0.00	0.00	59.00	7	kswapd0
19:29:43	0	94	0.00	45.00	0.00	0.00	45.00	7	kswapd0
19:29:44	0	94	0.00	36.00	0.00	0.00	36.00	0	kswapd0
19:29:45	0	94	0.00	27.00	0.00	1.00	27.00	0	kswapd0
19:29:46	0	94	0.00	26.00	0.00	0.00	26.00	2	kswapd0
19:29:47	0	94	0.00	57.00	0.00	0.00	57.00	7	kswapd0
19:29:48	0	94	0.00	41.00	0.00	0.00	41.00	7	kswapd0
19:29:49	0	94	0.00	38.00	0.00	0.00	38.00	7	kswapd0
19:29:50	0	94	0.00	47.00	0.00	0.00	47.00	7	kswapd0
19:29:51	0	94	0.00	43.00	0.00	1.00	43.00	7	kswapd0
19:29:52	0	94	0.00	36.00	0.00	0.00	36.00	7	kswapd0
19:29:53	0	94	0.00	39.00	0.00	0.00	39.00	2	kswapd0
19:29:54	0	94	0.00	46.00	0.00	0.00	46.00	7	kswapd0
19:29:55	0	94	0.00	43.00	0.00	0.00	43.00	7	kswapd0
19:29:56	0	94	0.00	39.00	0.00	0.00	39.00	7	kswapd0
19:29:57	0	94	0.00	29.00	0.00	1.00	29.00	1	kswapd0
19:29:58	0	94	0.00	17.00	0.00	0.00	17.00	4	kswapd0

	kbmemfree	kbavail	kbmemused	%memused	kbbuffers	kbcached	kbcommit
19:29:33	4464588	6742648	4420876	38.12	6156	2032600	179730872
19:29:34	4462572	6740784	4422752	38.13	6156	2032752	179739004
19:29:36	4381512	6740856	4422420	38.13	6156	2114144	179746508
19:29:37	3619456	6741840	4421588	38.12	6156	2877032	179746652
19:29:38	2848184	6740720	4422472	38.13	6164	3646188	179746652
19:29:39	2436336	6739452	4423720	38.14	6164	4056772	179746652
19:29:40	1712660	6737700	4425140	38.15	6164	4779020	179746604
19:29:41	810664	6738020	4425004	38.15	6164	5681152	179746604
19:29:42	673756	6779120	4373200	37.71	5656	5869928	179746604
19:29:43	688480	6782024	4371012	37.69	5648	5856940	179750048
19:29:44	688956	6789208	4364280	37.63	5584	5863272	179750048
19:29:45	740768	6804560	4348772	37.49	5524	5827248	179750000
19:29:46	697936	6810612	4342768	37.44	5524	5876048	179750048
19:29:47	734504	6818716	4334156	37.37	5512	5849188	179750000
19:29:48	771696	6828316	4324180	37.28	5504	5820948	179762260
19:29:49	691944	6838812	4313108	37.19	5476	5912444	179749952
19:29:50	679392	6844496	4306892	37.13	5452	5931356	179749952
19:29:51	768528	6868080	4284224	36.94	5412	5865704	176317452
19:29:52	717880	6893940	4259968	36.73	5400	5942368	176317404
19:29:53	712408	6902660	4251268	36.65	5372	5956584	176318376
19:29:54	707184	6917512	4236160	36.52	5344	5976944	176318568
19:29:55	703172	6921608	4232332	36.49	5292	5984836	176318568
19:29:56	73256	6933020	4220864	36.39	5212	5966340	176318568
19:29:57	723308	6936340	4217280	36.36	5120	5979816	176318568
19:29:58	732148	6942972	4210680	36.30	5108	5977656	176311064

Buffer IO 写

	UID	PID	%usr	%system	%guest	%wait	%CPU	CPU	Command
19:31:31	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:32	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:33	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:34	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:35	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:36	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:37	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:38	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:39	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:40	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:41	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:42	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:43	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:44	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:45	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:46	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:47	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:48	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0
19:31:49	0	94	0.00	0.00	0.00	0.00	0.00	4	kswapd0

	kbmemfree	kbavail	kbmemused	%memused	kbbuffers	kbcached	kbcommit
19:31:31	4816812	6928788	4225812	36.43	5148	1879678	176322636
19:31:32	4781880	6889428	4265592	36.78	5148	1874860	176322636
19:31:33	4758972	6822588	4332376	37.35	5148	1830984	176322636
19:31:34	4850248	6766480	4387840	37.83	5148	1684244	176322636
19:31:35	4864416	6741676	4413256	38.05	5148	1864900	176322636
19:31:36	4637900	6681480	4473436	38.57	5148	1810996	176322588
19:31:37	4502108	6595508	4559500	39.31	5148	1860724	176322492
19:31:38	4498844	6551068	4603928	39.69	5148	1819528	176322492
19:31:39	4498812	6587396	4567340	39.38	5148	1856116	176322492
19:31:40	4656784	6706252	4448372	38.35	5148	1817112	176322492
19:31:41	4635032	6673328	4481436	38.64	5148	1805816	176322492
19:31:42	4636852	6679736	4474884	38.58	5148	1810548	176322492
19:31:43	4654740	6669104	4485544	38.67	5148	1782000	176322444
19:31:44	4821604	6693156	4461848	38.47	5148	1638864	176322444
19:31:45	4707548	6728796	4426400	38.16	5148	1788368	176322444
19:31:46	4683996	6747632	4407348	38.00	5148	1830968	176322444
19:31:47	4694648	6773808	4381320	37.78	5148	1846376	176322624
19:31:48	4663784	6730212	4424776	38.15	5148	1833784	176322772

Uncached buffer IO 写

是否有更好方案解决 f2fs Uncached buffer I/O 写场景下性能下降问题？



Uncached buffer I/O 写为了能尽快释放 page cache 占用内存，其会在数据写入 page cache 后主动触发一次回写，这种方法类似于 sync 操作，进而会导致写性能降低。但是在 Jens Axboe 的测试环境，即磁盘性能非常好的环境下 UBI0 写性能会有提升，但是在绝大部分场景下磁盘的性能是低于内存的，故 UBI0 写会导致写性能降低。是否有更优方案来解决 UBI0 写性能降低的问题？

```
writing bs 65536, uncached 0
 1s: 196035MB/sec
 2s: 132308MB/sec
 3s: 132438MB/sec
 4s: 116528MB/sec
 5s: 103898MB/sec
 6s: 108893MB/sec
 7s: 99678MB/sec
 8s: 106545MB/sec
 9s: 106826MB/sec
10s: 101544MB/sec
11s: 111044MB/sec
12s: 124257MB/sec
13s: 116031MB/sec
14s: 114540MB/sec
15s: 115011MB/sec
16s: 115260MB/sec
17s: 116068MB/sec
18s: 116096MB/sec
```

```
writing bs 65536, uncached 1
 1s: 198974MB/sec
 2s: 189618MB/sec
 3s: 193601MB/sec
 4s: 188582MB/sec
 5s: 193487MB/sec
 6s: 188341MB/sec
 7s: 194325MB/sec
 8s: 188114MB/sec
 9s: 192740MB/sec
10s: 189206MB/sec
11s: 193442MB/sec
12s: 189659MB/sec
13s: 191732MB/sec
14s: 190701MB/sec
15s: 191789MB/sec
16s: 191259MB/sec
17s: 190613MB/sec
18s: 191951MB/sec
```

Jens Axboe 环境写性能：

<https://git.kernel.org/pub/scm/linux/kernel/git/akpm/mm.git/commit/?h=mm-new&id=d47c670061b5f9481ce494cd6c45078be301620e>

THANKS